

# CQS Summer Institute: Machine Learning and Statistics in Data Science

Matthew S. Shotwell, Ph.D.

Department of Biostatistics  
Vanderbilt University Medical Center  
Nashville, TN, USA


August 8, 2019

# Course Overview

- ▶ Syllabus and R code:
- ▶ <https://github.com/biostatmatt/cqs-ml-stat-r>
- ▶ Monday: Intro and Data Management
- ▶ Tuesday: Supervised Learning Part 1
- ▶ Wednesday: Supervised Learning Part 2
- ▶ Thursday: Supervised Learning Part 3
- ▶ Friday: Unsupervised Learning

# Deep learning

- ▶ Deep learning uses deep NNs
- ▶ Deep NNs are simply NNs with many layers, complex connectivity, and processing steps between layers:



# Complex NNs in R

- ▶ No (good) native R libraries for complex NNs
- ▶ R can interface to good libraries, notably Keras
- ▶ See <https://keras.rstudio.com/>
- ▶ “Deep Learning with R” by François Chollet and J. J. Allaire is a great resource, and provided a basis for this presentation, including the examples many of the figures.

# Keras in R

- ▶ Keras is a high-level model-building library for NNs; legos for NNs
- ▶ Keras can be used in R, Python, and other software
- ▶ Low-level computing handled by a “backend” library
- ▶ Backend libraries include TensorFlow, Theano, CNTK



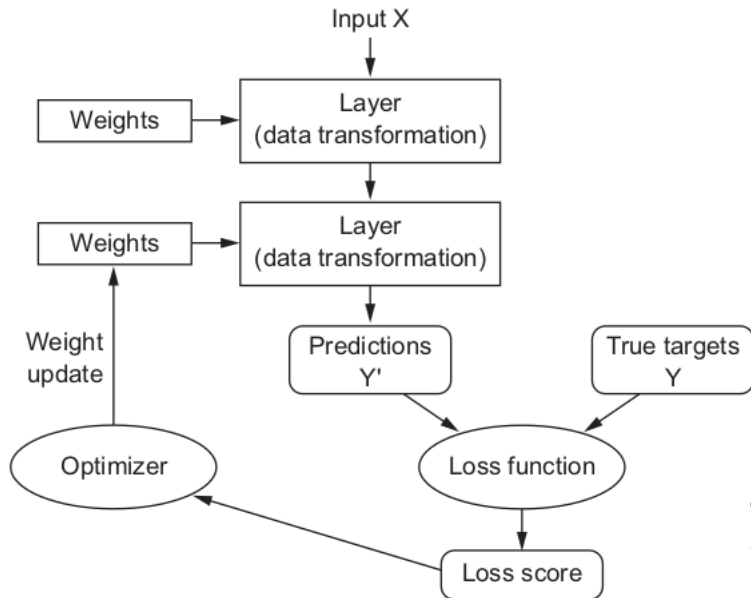
# Alternative method to set up Keras workstation

- ▶ set up a cloud computer on Amazon AWS
- ▶ get one that's already set up with Keras
- ▶ take advantage of GPU computing
- ▶ GPU gives 5 to 10 times speed improvement
- ▶ costs about \$0.90/hr
- ▶ [https://tensorflow.rstudio.com/tools/cloud\\_gpu](https://tensorflow.rstudio.com/tools/cloud_gpu)

# Modeling steps in Keras overview

1. define training data: input and target tensors
2. define network of layers mapping inputs to targets
3. choose loss function, optimizer and metrics (e.g., accuracy)
4. do optimization

# NN in Keras Overview






# Data representation in Keras

- ▶ in Keras, data are stored as tensors
- ▶ scalars - 0D tensors
- ▶ vectors - 1D tensors
- ▶ matrices - 2D tensors
- ▶ every tensor has attributes
  - ▶ rank, number of axes
  - ▶ shape, dimensions in each axis
  - ▶ type, data type, e.g., integer
- ▶  $100 \times 15$  numeric matrix has rank 2, shape (100, 5), and type 'float'
- ▶ in Keras, all data must usually be floating point numeric data

# 4D color image data tensor



## Step 1. Load and prep data in Keras

- ▶ Keras as some built-in data, including MNIST zipcode data
- ▶ first dimension is the 'sample axis'
- ▶ B&W image data have three axes: (samples, height, width)
- ▶ color image data typically have extra 'channel' axis

```
mnist <- dataset_mnist()  
train_images <- mnist$train$x  
train_labels <- mnist$train$y  
test_images <- mnist$test$x  
test_labels <- mnist$test$y  
  
> str(train_images)  
int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 0 ...  
> str(train_labels)  
int [1:60000(1d)] 5 0 4 1 9 2 1 3 1 4 ...
```

## Step 1. Load and prep data in Keras

- ▶ Keras as some built-in data, including MNIST zipcode data
- ▶ first dimension is the 'sample axis'
- ▶ B&W image data have three axes: (samples, height, width)
- ▶ color image data typically have extra 'channel' axis

```
train_images <- array_reshape(train_images, c(60000, 28 * 28))  
train_images <- train_images / 255  
test_images <- array_reshape(test_images, c(10000, 28 * 28))  
test_images <- test_images / 255  
  
train_labels <- to_categorical(train_labels)  
test_labels <- to_categorical(test_labels)
```

## Step 2. Defining Keras model

- ▶ the `keras_model_sequential()` function creates feed-forward NNs
- ▶ this type is made of “linear stack” of layers

```
model <- keras_model_sequential() %>%  
  layer_dense(  
    units = 32,  
    activation="relu",  
    input_shape = c(28*28)) %>%  
  layer_dense(  
    units = 10,  
    activation = "softmax")
```

# Choosing loss and last layer activation

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

## Step 3. Compile Keras model

- ▶ the `compile()` to specify optimizer, loss function, and metrics to monitor
- ▶ `compile()` modifies the model “in place”

```
model %>% compile(  
  optimizer = "rmsprop",  
  loss = "categorical_crossentropy",  
  metrics = c("accuracy")  
)
```

## Step 4. Fit Keras model

- use the `fit()` to pass the input and target data to the model for optimization, and to select the batch size and number of training epochs


```
model %>% fit(  
  train_images,  
  train_labels,  
  batch_size = 128,  
  epochs = 10  
)
```



# Keras example 1: Classifying movie reviews

- ▶ IMDB data set has 50k polarized reviews
- ▶ written reviews are a sequence of words preprocessed into a sequence of integers that index a dictionary
- ▶ human has read each review and classified each as positive or negative
- ▶ about 50% positive review, 50% negative
- ▶ create NN to classify reviews

# IMDB Classification Network




# Classifying movie reviews:

`keras-imdb-reviews.R`

# Convnets

- ▶ convnets (convolutional NNs)
- ▶ convolution involves weight sharing and local connectivity
- ▶ hidden layers are grouped into filters
- ▶ each filter is a “shape detector”
- ▶ multiple layers can learn more complex shapes

# Layers of shape detectors



# Convolution

A single filter does this:

# Convolution

For each convolutional layer:

- ▶ must specify size of patches from input
- ▶ must specify the number of filters
- ▶ must specify padding: “valid” or “same”
- ▶ must specify stride

# Padding: valid





Figure 5.5 Valid locations of  $3 \times 3$  patches in a  $5 \times 5$  input feature map



# Padding: same




Figure 5.6 Padding a  $5 \times 5$  input in order to extract 25  $3 \times 3$  patches

Stride:  $2 \times 2$





Figure 5.7  $3 \times 3$  convolution patches with  $2 \times 2$  strides

# **Classifying zipcodes:**

`keras-mnist-zipcodes.R`

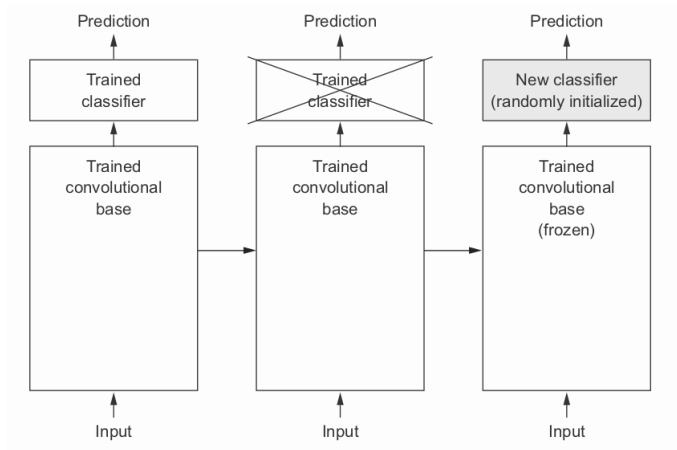
# Starting from scratch vs. pretrained convnet

- ▶ it's possible to train deep NN on few hundred images
- ▶ alternately, can fine-tuning a pretrained network
- ▶ or use pretrained network for feature extraction
- ▶ pretrained network with large general dataset can serve as a generic model for visual classification that we can use in other problems; this is a feature of deep learning that isn't available in shallow methods

# Using pretrained convnet: feature extraction/fine-tuning

- ▶ use pretrained network to extract features
- ▶ from “convolutional base” of pretrained network
- ▶ use derived features to train new convnet

# Using pretrained convnet: feature extraction/fine-tuning



# Using pretrained convnet: feature extraction/fine-tuning

- Get convolutional base from VGG16 NN trained on ImageNet

```
conv_base <- application_vgg16(  
  weights = "imagenet",  
  include_top = FALSE,  
  input_shape = c(150, 150, 3)  
)
```

# Using pretrained convnet: feature extraction/fine-tuning

```
> conv_base  
Model  
Model: "vgg16"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590880
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590880
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		



# Using pretrained convnet: feature extraction/fine-tuning

- Add the last couple of layers onto pretrained convolutional base

```
model <- keras_model_sequential() %>%  
  conv_base %>%  
  layer_flatten() %>%  
  layer_dense(units = 256, activation = "relu") %>%  
  layer_dense(units = 1, activation = "sigmoid")  
  
freeze_weights(conv_base)
```

# Using pretrained convnet: feature extraction/fine-tuning

```
> model
```

```
Model
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_8 (Dense)	(None, 256)	2097408
dense_9 (Dense)	(None, 1)	257
Total params: 16,812,353		
Trainable params: 2,097,665		
Non-trainable params: 14,714,688		

# Feature extraction examples

- ▶ <https://github.com/jjallaire/deep-learning-with-r-notebooks/tree/master/notebooks>
- ▶ 5.2 and 5.3